

A Workflow Support System for Computer-Aided Tuberculosis Screening from Chest X-Rays

Raphael B. Alampay , Irish Danielle Morales , Ruben A. Saulog  and Patricia Angela R. Abu 
Department of Information Systems and Computer Science, Ateneo de Manila University, Quezon City, Philippines

Abstract—Tuberculosis (TB) causes more deaths than any other infectious disease worldwide. While computer-aided detection (CAD) systems used in TB screening programs achieve high diagnostic accuracy, several practical challenges emerge in the screening workflow. Commercial TB CAD products often do not account for X-ray image quality, post-deployment monitoring, and explanation of findings in natural language for screening teams, which often do not involve a radiologist. We present a prototype workflow support system for TB screening to address these issues. Its lead contribution is a clinically grounded LLM explanation module that describes model findings in natural language, a feature not yet present in the commercial products we reviewed. Supporting modules add an input quality gate to assess X-ray image quality and a post-deployment monitoring dashboard. We also develop and benchmark custom models for TB classification and localization. Our results show that ... (TODO).

Keywords—Tuberculosis; Computer-Aided Detection; Chest X-ray; Medical Imaging; Computer Vision

I. INTRODUCTION

Tuberculosis (TB) causes more deaths than any other infectious disease worldwide [1]. It is caused by airborne *Mycobacterium tuberculosis* and primarily affects the lungs [3]. Low- and middle-income countries bear a disproportionate share of the burden, compounded by limited access to diagnostics, delayed presentation, and overburdened health infrastructure [2].

TB computer-aided detection (CAD) systems have matured into routine clinical use. The World Health Organization (WHO) endorsed CAD systems as an alternative to human readers for people aged 15 years and older [6]. TB screening programs often use mobile X-ray units and CAD systems, where CAD substitutes for human readers [25] [26]. In 2025, WHO approved six commercial products as meeting its TB screening performance standards [7]. In one reader study, AI-CAD assistance improved radiologist performance and enabled non-radiologists to achieve similar performance as radiologists [9]. Several reviews report model accuracy approaching WHO targets across populations [8], [39], [40], indicating that the field has moved beyond improving accuracy toward real-world use in high-volume screening programs.

A. Gaps in Screening Workflow Integration

Reviews and deployment studies consistently identify gaps in practical implementation of CAD systems. Among CXR AI studies, workflow integration, local validation, and human factors remain under-studied [10], [8]. An India deployment identified barriers in workflow, connectivity, and interpretation [36], a Vietnam deployment showed only 47.3% agreement between physician and CAD interpretations at three

months [37], and clinicians across countries valued CAD as a complement rather than a replacement [38]. The following gaps emerge in the literature:

Image quality varies in real-world use. Chest X-ray (CXR) is the most widely used modality for TB screening [4]. TB CAD systems are trained on curated datasets, but CXR images may be poorly positioned, clipped, underexposed, or photographed from analog film [27], especially in community screening programs using mobile X-ray units. Poor input quality degrades performance, yet most systems do not expose user-facing quality assessment before inference [10], [11].

Heatmaps for explaining model findings lack clinical utility. All commercial TB CAD products in Table I provide heatmaps for model explainability. However, heatmaps are less intuitive for radiologists than discrete localized findings [42], often overlap poorly with expert-identified regions, and do not convey the contextual reasoning used during diagnosis [28]. A TB CAD acceptability study found that users wanted to understand why the model reached its conclusions [46]. Textual descriptions and case-based reasoning have been proposed as more clinically meaningful alternatives [29], [12].

Post-deployment monitoring is often lacking. Deployed performance may shift due to population, equipment, or software changes [53]. The U.S. Food and Drug Administration has identified these as risks to AI-enabled medical devices and initiated research on proactive post-market monitoring [32], but in practice this remains an operational gap [10].

B. Related Work

Related literature has explored models to explain findings in natural language, although this is not yet implemented in deployed systems or commercial products. In a research prototype, Siemens Healthineers paired a combined classification-localization model [13] with a domain-adapted LLM to generate CXR findings [12]. General medical vision-language models such as LLaVA-Med [14] and CXR-LLaVA [30] demonstrate image-conditioned radiology text generation, and recent surveys catalogue the rapid growth of vision-language models for medical report generation [31].

We compare against the six WHO-approved TB CAD products in Table I: qXR v4.0 [19], CAD4TB v7 [20], Lunit INSIGHT CXR v3.1 [21], DrAid v2.4.4–6 [22], Genki Edge v3.4.2 [23], and InferRead DR Chest v1.0.1.1 [24] [7]. While all provide heatmap localization, post-deployment monitoring is uncommon, input quality assessment is largely absent, and no product is known to explain findings in natural language.

TABLE I. FEATURE COMPARISON WITH COMMERCIAL TB CAD PRODUCTS

Feature	qXR [19]	CAD4TB+ [20]	Lunit [21]	DrAid [22]	Genki [23]	InferRead [24]	Proposed
Chest X-ray quality assessment				n/a	n/a		✓
Monitoring and analytics dashboard			✓	n/a	n/a		✓
Heatmap/overlay of localization findings	✓	✓	✓	✓	✓	✓	✓
Explanation of findings in natural language				n/a	n/a		✓

✓ marks a feature documented as present; blank marks a feature documented as absent or not offered; n/a marks features for which no documentation was found. The monitoring/analytics dashboard column denotes a post-deployment dashboard for tracking model usage and performance over time.

C. Scope and Contribution

This paper presents a workflow support system for TB screening that targets the gaps above. Its contributions are:

- 1) A **natural language explanation module** (lead contribution) that describes detected abnormalities with clinical grounding, supporting bounding box-level visual evidence.
- 2) An **input quality assessment module** that screens CXR images for acquisition defects before inference.
- 3) A **post-deployment monitoring dashboard** tracking system behavior and performance over time.

As a secondary contribution, we develop and benchmark custom models for classification (Section IV) and localization (Section V).

II. SYSTEM OVERVIEW

Figure 1 describes the system overview. A CXR (DICOM or standard image format) is ingested from the screening unit's X-ray or uploaded directly, then evaluated for acquisition quality. Images passing quality screening are scored by the TB classification module. For images flagged as TB, localization (object detection) is done to detect active and latent TB regions. Findings are passed to a module that generates a natural language description of the findings based in clinical context. Outputs are displayed in the interface, and each case is logged for post-deployment monitoring.

III. INPUT QUALITY ASSESSMENT

IV. TUBERCULOSIS CLASSIFICATION

The TB classification module flags each CXR as either healthy, sick non-TB, or TB as seen in Figure 1.

A. Dataset

Classification and localization models are trained on TBX11K [48], a public CXR dataset with image-level labels (*healthy*, *sick non-TB*, *TB*) and bounding-box-level labels (*active-TB*, *latent-TB*). We use a simplified redistribution¹ with 8,400 CXRs (3,800 healthy, 3,800 sick non-TB, 800 TB), with a 79%/21% train/validation split (6,600 and 1,800 images). TBX11K is selected since it is the largest public TB CXR dataset combining a TB-specific sample large enough to train from scratch, expert bounding-box annotations, a sick non-TB class, and a demographic and class composition that reflect real-world clinical settings [49].

¹<https://www.kaggle.com/datasets/vbookshelf/tbx11k-simplified>

The official TBX11K test set is withheld for an online challenge, so we evaluate on the public validation set. Our results are therefore not directly comparable to SymFormer [49] or other challenge entries on the private test set.

B. Data Preprocessing

Classification models are trained without preprocessing or data augmentation. CXRs unfit for inference are already filtered out by the input quality assessment module. CXRs are also acquired under a standardized posteroanterior protocol and are largely uniform in pose, so aggressive augmentation risks introducing distribution shift.

C. Model Architecture

We evaluate nine classifiers: six baselines (EfficientNet-B0 [60], MobileNetV3-Large [59], DenseNet-121 [58], ResNet-18, ResNet-50 [57], ConvNeXt-Tiny [61]) and three custom models (DraxNet, Drax-MobileNetV3-Large, FlipR).

1) *DraxNet*: DraxNet (17.0M parameters) modifies a ResNet-18 backbone so the fourth layer consists of two ResNet blocks, each with a custom mixer (*DraxBlock*) inserted between its two convolutions. All other components are unchanged. *DraxBlock* consists of a ConvNeXt local convolutional branch followed by an optional self-attention branch, fused via stochastic-depth residuals. DraxNet, short for Dynamic Residual Attention eXchange, reserves the heavier mixer for the final, lowest-resolution stage, adding attention capacity with fewer parameters than a full attention backbone.

2) *Drax-MobileNetV3-Large*: Drax-MobileNetV3-Large (6.1M parameters) applies the DraxBlock to a parameter-efficient backbone. The MobileNetV3-Large feature extractor and head are unchanged, and a bottlenecked *DraxBlock* is inserted after the final feature stage. This tests whether the DraxBlock generalizes beyond a ResNet-style backbone.

3) *FlipR*: FlipR (2.8M parameters) inserts a *FlipR* block between layers of an otherwise unchanged ResNet-18. The *FlipR* block computes the difference between each feature map and its horizontally flipped counterpart, smoothed by average pooling to account for imperfect anatomical symmetry, then broadcasts a gated mask across all channels. The block is inspired by SymAttention [49], which observes that bilateral asymmetry in CXRs is a key indicator of TB. FlipR tests whether a lightweight symmetry gate can match heavier attention approaches.

D. Model Training

All models are trained from scratch with a fixed random seed, batch size 16, 512×512 input, 100 epochs, Adam at




Fig. 1. Overview of the proposed TB CAD pipeline.

a fixed learning rate of 10^{-3} , no schedule, and unweighted cross-entropy loss. No hyperparameter tuning was done.

E. Results

Classification results are reported in Table II.

MobileNetV3-Large is the top classifier, while FlipR achieves comparable results with half the parameters. Given the small ($\pm 1\%$) differences, ordering among the top four (FlipR, EfficientNet, MobileNetV3-Large, Drax-MobileNetV3-Large) is likely due to stochasticity. They differ most on sensitivity ($\pm 2.5\%$), so we interpret MobileNetV3-Large (96.67% sensitivity) as the top model, since sensitivity is the clinically dominant metric in TB screening. FlipR reaches comparable results on other metrics at roughly half the size (2.8M vs 5.5M), supporting the SymFormer [49] hypothesis that bilateral asymmetry is highly indicative of TB.

All models meet WHO targets for TB triage. Sensitivity ranges from 92.50 to 96.67%, clearing the WHO target of 90% sensitivity. Specificity ranges from 97.98 to 99.91%, well above the WHO target of 70% specificity [2]. Sensitivity is more clinically important than specificity in screening, since a missed TB case risks untreated transmission while a false positive triggers only confirmatory testing. The top model, MobileNetV3-Large, misses only 3.3% of TB cases.

The high-specificity, lower-sensitivity profile is likely attributable to TBX11K rather than model architecture. Models in the literature consistently show the same pattern on TBX11K. The dataset is constructed to reward high specificity, since the common failure mode on other TB datasets is over-predicting TB and collapsing specificity to near-zero [49].

Lightweight models are preferred for TB screening on both performance and deployment grounds. Smaller models outperform larger ones across the board: MobileNetV3-Large (5.5M), EfficientNet-B0 (5.3M), and FlipR (2.8M) beat all larger models in accuracy and F1. Due to a modest TBX11K training set, a minority TB class, and no class-imbalance handling, larger networks appear to overfit the majority healthy and sick non-TB classes rather than learn discriminative TB features. TB screening also runs in resource-limited settings often without a dedicated GPU, so a lightweight model is preferable in deployment.

V. TUBERCULOSIS LOCALIZATION

A. Dataset

Localization uses the TBX11K bounding-box annotations, with each region labeled active TB or latent TB [48]. Active TB denotes ongoing, transmissible infection requiring treatment; latent TB denotes prior, inactive infection that is not contagious but may reactivate. As Figure 1 shows, active TB is more visually salient and latent TB is more subtle. A TB-labeled CXR may contain only active-TB lesions (924 images), only latent-TB lesions (212 images), or both (54 images), a composition consistent with real-world conditions.

B. Model Architecture

We evaluate YOLO26 and a custom variant in which the detector's default feature extractor is replaced by the DraxNet backbone used in classification. The detection head and training protocol are unchanged, isolating the backbone's effect on localization.

TABLE II. TUBERCULOSIS CLASSIFICATION RESULTS

Model	AUC _{TB}	Acc. (%)	Sens. (%)	Spec. (%)	Avg. Prec. (%)	Avg. Rec. (%)	Avg. F1 (%)	Params
FlipR	0.9984	98.97	94.17	99.91	99.01	97.70	98.34	2.8M
EfficientNet-B0	0.9995	98.89	95.00	99.65	98.29	97.87	98.07	5.3M
MobileNetV3-Large	0.9995	99.21	96.67	99.82	98.97	98.54	98.75	5.5M
Drax-MobileNetV3-Large	0.9989	98.81	95.00	99.65	98.23	97.81	98.02	6.1M
DenseNet-121	0.9975	98.57	92.50	99.56	97.81	96.97	97.38	8.0M
ResNet-18	0.9982	98.17	93.33	99.21	96.71	96.90	96.80	11.7M
DraxNet	0.9988	98.41	95.00	99.47	97.51	97.51	97.51	17.0M
ResNet-50	0.9978	97.70	92.50	99.47	96.95	96.33	96.63	25.6M
ConvNeXt-Tiny	0.9952	97.14	95.00	97.98	93.63	96.58	94.97	28.6M

AUC_{TB} is the area under the ROC curve for the TB class. Sensitivity is recall for the TB class. Specificity is recall for the non-TB class (healthy and sick non-TB combined). Precision, recall, and F1 are averaged across the three classes (healthy, sick non-TB, TB).




Fig. 2. Localization precision-recall (left) and F1-confidence (right) curves for YOLO26 (top) and YOLO26 with the DraxNet backbone (bottom). Both sustain high precision while recall saturates near 0.26, and the latent-TB curve collapses at low confidence, so localization functions as a high-precision rule-in cue rather than a standalone detector.

C. Model Training

Localization uses the same training configuration as the classification models (Section IV).

TABLE III. TUBERCULOSIS LOCALIZATION RESULTS

Model	Class	mAP ₅₀	mAP ₅₀₋₉₅
YOLO26	ATB	—	0.2110
YOLO26	LTB	—	0.0410
YOLO26 + DraxNet	ATB	—	0.2748
YOLO26 + DraxNet	LTB	—	0.0310

ATB = active TB; LTB = latent TB.

D. Results

Localization results per class are reported in Table III.

Active TB is more likely to be detected than latent TB.

The very low mAP₅₀₋₉₅ for latent TB reflects the TBX11K class imbalance (active-TB regions outnumber latent-TB regions roughly 4:1) and the subtle appearance of latent-TB lesions. The TBX11K reference reports the same pattern across all detection methods, attributing it to imbalance [49], and it is consistent with human readers distinguishing latent TB less reliably [5]. The bias toward active TB also aligns with clinical triage priorities, since active TB requires urgent treatment.

The DraxNet backbone specializes toward active TB. YOLO26 with the DraxNet backbone raises active-TB mAP₅₀₋₉₅ from 0.2110 to 0.2748 (+0.0638) while slightly reducing latent-TB (0.0410 to 0.0310). The gain is driven almost entirely by precision (0.74 → 0.82). Recall is unchanged at roughly 0.26.

Recall, not precision, is the bottleneck. Both detectors localize confidently but find only about a quarter of the annotated TB regions. Figure 2 shows precision staying high while recall saturates near 0.26, and the latent-TB curve collapsing at low confidence for both detectors.

Localization is useful to confirm TB, not to rule it out. The high-precision, low-recall profile dictates clinical use: a flagged lesion is likely correct and warrants focused radiologist review, but the detector must not exclude TB since most annotated lesions are missed. Localization is therefore most reliable as supporting region-level evidence for the classifier and LLM explanation modules, rather than as a standalone detector.

VI. MULTIMODAL EXPLANATION

VII. STRUCTURED EVIDENCE INTERFACE

VIII. POST-DEPLOYMENT MONITORING

IX. DISCUSSION

A. Limitations

B. Future Work

X. CONCLUSION

DECLARATION ON GENERATIVE AI

Generative AI tools were used to assist in drafting and preparing this manuscript. The conception of the research, methodology, and conclusions are the work of the authors.

REFERENCES

- [1] World Health Organization, "Global Tuberculosis Report 2024," Geneva: WHO, 2024. [Online]. Available: <https://www.who.int/teams/global-tuberculosis-programme/tb-reports/global-tuberculosis-report-2024>
- [2] World Health Organization, "WHO consolidated guidelines on tuberculosis. Module 2: Screening – Systematic screening for tuberculosis disease," Geneva: WHO, 2021. [Online]. Available: <https://www.who.int/publications/i/item/9789240022676>
- [3] M. Pai, M. A. Behr, D. Dowdy, K. Dheda, M. Divangahi, C. C. Boehme, A. Ginsberg, S. Swaminathan, M. Spigelman, H. Getahun, D. Menzies, and M. Raviglione, "Tuberculosis," *Nature Reviews Disease Primers*, vol. 2, art. 16076, 2016, doi: 10.1038/nrdp.2016.76.
- [4] J. Burrill, C. J. Williams, G. Bain, G. Conder, A. L. Hine, and R. R. Misra, "Tuberculosis: A radiologic review," *RadioGraphics*, vol. 27, no. 5, pp. 1255–1273, 2007, doi: 10.1148/rg.275065176.
- [5] Y. R. Choi, S. H. Yoon, J. Kim, J. Y. Yoo, H. Kim, and K. N. Jin, "Chest radiography of tuberculosis: determination of activity using deep learning algorithm," *Tuberculosis and Respiratory Diseases*, vol. 86, no. 3, pp. 226–233, 2023, doi: 10.4046/trd.2023.0020.
- [6] World Health Organization, "WHO operational handbook on tuberculosis. Module 2: Screening," Geneva: WHO, 2021. [Online]. Available: <https://www.who.int/publications/i/item/9789240110373>
- [7] World Health Organization, "Use of computer-aided detection software for tuberculosis screening: WHO policy statement," Geneva: WHO, 2025. [Online]. Available: <https://www.who.int/publications/i/item/9789240110373>
- [8] H. K. Ahmad *et al.*, "Machine learning augmented interpretation of chest X-rays: A systematic review," *Diagnostics*, vol. 13, no. 4, art. 743, 2023. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC9955112/>
- [9] P. G. Anderson *et al.*, "Deep learning improves physician accuracy in the comprehensive detection of abnormalities on chest X-rays," *Scientific Reports*, vol. 14, 2024. [Online]. Available: <https://www.nature.com/articles/s41598-024-76608-2>
- [10] J. Mongan *et al.*, "Artificial intelligence and machine learning for chest radiography," *Radiology*, vol. 307, no. 3, 2023. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC10128969/>
- [11] A. J. DeGrave, J. D. Janizek, and S.-I. Lee, "AI for radiographic COVID-19 detection selects shortcuts over signal," *Nature Machine Intelligence*, vol. 3, pp. 610–619, 2021. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC10047562/>
- [12] M. D. Danu, G. Marica, S. K. Karn, B. Georgescu, A. Mansoor, F. Ghesu, L. M. Itu, C. Suciu, S. Grbic, O. Farri, and D. Comaniciu, "Generation of radiology findings in chest X-ray by leveraging collaborative knowledge," *Procedia Computer Science (ITQM 2023)*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.10448>
- [13] F. C. Ghesu, B. Georgescu, A. Mansoor, Y. Yoo, D. Neumann, P. Patel, R. S. Vishwanath, J. M. Balter, Y. Cao, S. Grbic, and D. Comaniciu, "Self-supervised learning from 100 million medical images," *arXiv preprint arXiv:2201.01283*, 2022. [Online]. Available: <https://arxiv.org/abs/2201.01283>
- [14] C. Li *et al.*, "LLaVA-Med: Training a large language-and-vision assistant for biomedicine in one day," *Advances in Neural Information Processing Systems*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.00890>
- [15] P. Lakhani and B. Sundaram, "Deep learning at chest radiography: Automated classification of pulmonary tuberculosis by using convolutional neural networks," *Radiology*, vol. 284, no. 2, pp. 574–582, 2017.
- [16] F. Pasa *et al.*, "Efficient deep network architectures for fast chest X-ray tuberculosis screening and visualization," *Scientific Reports*, vol. 9, 2019.
- [17] WHO Kobe Centre, "Determining the local calibration of computer-assisted detection (CAD) thresholds and other parameters." [Online]. Available: <https://wkc.who.int/resources/publications/i/item/determining-the-local-calibration-of-computer-assisted-detection-cad-thresholds-and-other-parameters>
- [18] A. Marquez *et al.*, "Performance of chest X-ray with computer-aided detection as a triage test for pulmonary TB in the Philippines," *BMC Global and Public Health*, 2025. [Online]. Available: <https://bmcbglobalpublichealth.biomedcentral.com/articles/10.1186/s44263-025-00198-y>
- [19] Qure.ai, "qXR documentation." [Online]. Available: <https://documentation.qure.ai/>
- [20] Delft Imaging, "CAD4TB+." [Online]. Available: <https://delft.care/cad4tb/>
- [21] Lunit, "INSIGHT CXR." [Online]. Available: <https://www.lunit.io/en/products/cxr>
- [22] VinBrain, "DrAid for TB screening." [Online]. Available: <https://draid.ai/>
- [23] DeepTek, "Genki: AI for tuberculosis screening." [Online]. Available: <https://www.deeptek.ai/genki>
- [24] Infervision, "InferRead DR Chest: instruction for use," (hosted on the Stop TB Partnership Digital Health Technology Hub). [Online]. Available: <https://stoptb.org/assets/documents/dhthub/User%20Manual%20Software%20-%20InferRead%20DR%20Chest.pdf>
- [25] Stop TB Partnership, "Screening and triage for TB using computer-aided detection (CAD) technology and ultra-portable X-ray systems: A practical guide." [Online]. Available: https://www.stoptb.org/sites/default/files/imported/page/Screening_and_Triage_for_TB_using_Computer-Aided_Detection_CAD_Technology_and_Ultra-portable_X-Ray_Systems-A_Practical_Guide_.pdf
- [26] FIND, "Digital chest radiography and computer-aided detection (CAD) solutions for tuberculosis diagnostics: Technology landscape analysis," 2021. [Online]. Available: https://www.finddx.org/wp-content/uploads/2023/02/20210401_technology_landscape_computer_aided_tb_FV_EN.pdf
- [27] Z. Desita, T. Tadesse, A. S. Bohlbro *et al.*, "Computer-aided analysis of photographed chest X-ray films performs well compared with trained radiologists," *Mayo Clinic Proceedings: Digital Health*, 2026. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC12955127/>
- [28] H. Chen, C. Gomez, C.-M. Huang, and M. Unberath, "Explainable medical imaging AI needs human-centered design: guidelines and evidence from a systematic review," *npj Digital Medicine*, vol. 5, no. 156, 2022. [Online]. Available: <https://www.nature.com/articles/s41746-022-00699-2>
- [29] K. Borys, Y. A. Schmitt, M. Nauta, C. Seifert, N. Krämer, C. M. Friedrich, and F. Nensa, "Explainable AI in medical imaging: An overview for clinical practitioners – Beyond saliency-based XAI approaches," *European Journal of Radiology*, vol. 162, art. 110786, 2023. [Online]. Available: <https://doi.org/10.1016/j.ejrad.2023.110786>
- [30] S. Lee, J. Youn, H. Kim, M. Kim, and S. H. Yoon, "CXR-LLaVA: a multimodal large language model for interpreting chest X-ray images," *European Radiology*, vol. 35, 2025. [Online]. Available: <https://link.springer.com/article/10.1007/s00330-024-11339-6>
- [31] I. Hartsock and G. Rasool, "Vision-language models for medical report generation and visual question answering: a review," *Frontiers in Artificial Intelligence*, vol. 7, art. 1430984, 2024. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC11611889/>
- [32] U.S. Food and Drug Administration, "Methods and tools for effective postmarket monitoring of artificial intelligence (AI)-enabled medical devices," 2024. [Online]. Available: <https://www.fda.gov/medical-devices/medical-device-regulatory-science-research-programs-conducted-osel/methods-and-tools-effective-postmarket-monitoring-artificial-intelligence-ai-enabled-medical-devices>
- [33] J. Rudolph *et al.*, "Post-deployment performance of a deep learning algorithm for normal and abnormal chest X-ray classification: A study at visa screening centers in the United Arab Emirates," *PLOS ONE*, 2024. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC11539241/>
- [34] A. S. Tejani, H. Elhalawani, L. Moy, M. Kohli, and C. E. Kahn Jr, "Artificial intelligence and radiology education," *Radiology: Artificial Intelligence*, vol. 5, no. 1, e220084, 2023. [Online]. Available: <https://pubs.rsna.org/doi/full/10.1148/ryai.220084>
- [35] M. A. Mazurowski, "Artificial intelligence for precision education in radiology," *British Journal of Radiology*, vol. 92, no. 1103, 2019. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC6849670/>
- [36] S. Vijayan *et al.*, "Implementing a chest X-ray artificial intelligence

- tool to enhance tuberculosis screening in India: Lessons learned,” *PLOS Digital Health*, vol. 2, no. 12, e0000404, 2023. [Online]. Available: <https://journals.plos.org/digitalhealth/article?id=10.1371/journal.pdig.0000404>
- [37] A. L. Innes *et al.*, “Computer-aided detection for chest radiography to improve the quality of tuberculosis diagnosis in Vietnam’s district health facilities: An implementation study,” *Tropical Medicine and Infectious Disease*, vol. 8, no. 11, art. 488, 2023. [Online]. Available: <https://doi.org/10.3390/tropicalmed8110488>
- [38] J. Creswell *et al.*, “Early user perspectives on using computer-aided detection software for interpreting chest X-ray images to enhance access and quality of care for persons with tuberculosis,” *BMC Global and Public Health*, vol. 1, no. 30, 2023. [Online]. Available: <https://bmcbglobalpublichealth.biomedcentral.com/articles/10.1186/s44263-023-00033-2>
- [39] A. J. Scott *et al.*, “Diagnostic accuracy of computer-aided detection during active case finding for pulmonary tuberculosis in Africa: A systematic review and meta-analysis,” *Open Forum Infectious Diseases*, vol. 11, no. 2, ofae020, 2024. [Online]. Available: <https://doi.org/10.1093/ofid/ofae020>
- [40] F. A. Khan *et al.*, “A systematic review and meta-analysis of artificial intelligence software for tuberculosis diagnosis using chest X-ray imaging,” 2025. [Online]. Available: <https://pubmed.ncbi.nlm.nih.gov/40529749/>
- [41] “Domain shift analysis in chest radiographs classification in a Veterans Healthcare Administration population,” *Journal of Imaging Informatics in Medicine*, 2025. [Online]. Available: <https://link.springer.com/article/10.1007/s10278-025-01494-7>
- [42] I. F. Ihongbe, S. Fouad, T. F. Mahmoud, A. Rajasekaran, and A. Bhatia, “Evaluating explainable artificial intelligence (XAI) techniques in chest radiology imaging through a human-centered lens,” *PLOS ONE*, vol. 19, no. 10, e0308758, 2024. [Online]. Available: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0308758>
- [43] S. Gaube *et al.*, “Care to explain? AI explanation types differentially impact chest radiograph diagnostic performance and physician trust in AI,” *Radiology*, vol. 313, no. 1, e233261, 2024. [Online]. Available: <https://pubs.rsna.org/doi/10.1148/radiol.233261>
- [44] S. Tonekaboni, S. Joshi, M. D. McCradden, and A. Goldenberg, “What clinicians want: Contextualizing explainable machine learning for clinical end use,” in *Proc. Machine Learning for Healthcare (MLHC)*, 2019. [Online]. Available: <https://arxiv.org/abs/1905.05134>
- [45] C. J. Cai, S. Winter, D. Steiner, L. Wilcox, and M. Terry, “‘Hello AI’: Uncovering the onboarding needs of medical practitioners for human-AI collaborative decision-making,” *Proc. ACM Human-Computer Interaction (CSCW)*, vol. 3, art. 104, 2019. [Online]. Available: <https://dl.acm.org/doi/10.1145/3359206>
- [46] D. Hua *et al.*, “Towards human-AI collaboration in radiology: A multidimensional evaluation of the acceptability of AI for chest radiograph analysis in supporting pulmonary tuberculosis diagnosis,” *JAMIA Open*, vol. 8, no. 1, ooae151, 2025. [Online]. Available: <https://academic.oup.com/jamiaopen/article/8/1/ooae151/8002370>
- [47] G. Chassagnon *et al.*, “Learning from the machine: AI assistance is not an effective learning tool for resident education in chest X-ray interpretation,” *European Radiology*, 2023. [Online]. Available: <https://pubmed.ncbi.nlm.nih.gov/37572190/>
- [48] Y. Liu, Y.-H. Wu, Y. Ban, H. Wang, and M.-M. Cheng, “Rethinking computer-aided tuberculosis diagnosis,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 2646–2655. [Online]. Available: <https://mmcheng.net/tb/>
- [49] Y. Liu, Y.-H. Wu, S.-C. Zhang, L. Liu, M. Wu, and M.-M. Cheng, “Revisiting computer-aided tuberculosis diagnosis,” *IEEE Trans. Pattern Analysis and Machine Intelligence*, 2023. [Online]. Available: <https://arxiv.org/abs/2307.02848>
- [50] F. Yang, P.-X. Lu, M. Deng *et al.*, “Annotations of lung abnormalities in the Shenzhen chest X-ray dataset for computer-aided screening of pulmonary diseases,” *Data*, vol. 7, no. 7, art. 95, 2022. [Online]. Available: <https://www.mdpi.com/2306-5729/7/7/95>
- [51] J. P. Cohen, J. D. Viviano, P. Bertin *et al.*, “TorchXRyVision: A library of chest X-ray datasets and models,” in *Proc. Medical Imaging with Deep Learning (MIDL)*, 2022. [Online]. Available: <https://arxiv.org/abs/2111.00595>
- [52] A. Bustos, A. Pertusa, J.-M. Salinas, and M. de la Iglesia-Vayá, “PadChest: A large chest X-ray image dataset with multi-label annotated reports,” *Medical Image Analysis*, vol. 66, art. 101797, 2020. [Online]. Available: <https://arxiv.org/abs/1901.07441>
- [53] S. G. Finlayson *et al.*, “The clinician and dataset shift in artificial intelligence,” *New England Journal of Medicine*, vol. 385, pp. 283–286, 2021. [Online]. Available: <https://pubmed.ncbi.nlm.nih.gov/34260843/>
- [54] J. Feng, R. V. Phillips, I. Malenica *et al.*, “Clinical artificial intelligence quality improvement: Towards continual monitoring and updating of AI algorithms in healthcare,” *npj Digital Medicine*, vol. 5, art. 66, 2022. [Online]. Available: <https://www.nature.com/articles/s41746-022-00611-y>
- [55] J. Merkow *et al.*, “Empirical data drift detection experiments on real-world medical imaging data,” *Nature Communications*, vol. 15, art. 1689, 2024. [Online]. Available: <https://www.nature.com/articles/s41467-024-46142-w>
- [56] A. Ghosh *et al.*, “Scalable drift monitoring in medical imaging AI,” *arXiv preprint arXiv:2410.13174*, 2024. [Online]. Available: <https://arxiv.org/abs/2410.13174>
- [57] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2016, pp. 770–778. [Online]. Available: <https://arxiv.org/abs/1512.03385>
- [58] G. Huang, Z. Liu, L. van der Maaten, and K. Q. Weinberger, “Densely connected convolutional networks,” in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 4700–4708. [Online]. Available: <https://arxiv.org/abs/1608.06993>
- [59] A. Howard *et al.*, “Searching for MobileNetV3,” in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2019, pp. 1314–1324. [Online]. Available: <https://arxiv.org/abs/1905.02244>
- [60] M. Tan and Q. V. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2019, pp. 6105–6114. [Online]. Available: <https://arxiv.org/abs/1905.11946>
- [61] Z. Liu, H. Mao, C.-Y. Wu, C. Feichtenhofer, T. Darrell, and S. Xie, “A ConvNet for the 2020s,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 11976–11986. [Online]. Available: <https://arxiv.org/abs/2201.03545>
- [62] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft COCO: Common objects in context,” in *Proc. European Conf. Computer Vision (ECCV)*, 2014, pp. 740–755. [Online]. Available: <https://arxiv.org/abs/1405.0312>